

Neural Networks

1 Basic Optimizer

In this course we will implement advanced optimization schemes, but in the first exercise we start with the basic Stochastic Gradient Descent (SGD).

Task:

Implement the class **Sgd** in the file “Optimizers.py” in folder “Optimization”.

- The **Sgd** constructor receives the **learning_rate** with data type float.
- Implement the method **calculate_update(weight_tensor, gradient_tensor)** that returns the updated weights according to the basic gradient descent update scheme.

2 Base Layer

We will realize a small layer oriented Deep Learning framework in this exercise. Layer oriented frameworks represent a higher level of abstraction to their users than graph oriented frameworks. This approach limits flexibility but enables easy experimentation using conventional architectures. Every layer in these architectures has to implement two fundamental operations: **forward(input_tensor)**, **backward(error_tensor)**. These operations are the basic steps executed during training and testing.

We distinguish between **trainable** and **non-trainable** layers. Trainable layers have parameters that are optimized during training (e.g. the Fully Connected Layer, which must be implemented in this task), while non-trainable layers remain fixed (e.g. the ReLU activation function).

Task:

Implement a class **BaseLayer** in the file “Base.py” in folder “Layers”.

- This class will be inherited by every layer in our framework. For information on inheritance in python, please refer to [here](#).
- Write a constructor for this class receiving no arguments. In this constructor, initialize a boolean member **trainable** with **False**. This member will be used to distinguish trainable from non-trainable layers.
- **Optionally**, you can add other members like a default **weights** parameter, which might come in handy.

3 Fully Connected Layer

The Fully Connected (FC) layer is the theoretic backbone of layer oriented architectures. It performs a linear operation on its input.

Task:

Implement a class **FullyConnected** in the file “FullyConnected.py” in folder “Layers”, that inherits the base layer that we implemented earlier. This class has to provide the methods **forward(input_tensor)** and **backward(error_tensor)** as well as the property **optimizer**.

- Write a constructor for this class, receiving the arguments (**input_size, output_size**). First, call its super-constructor. Set the inherited member **trainable** to **True**, as this layer has trainable parameters. Initialize the weights of this layer uniformly random in the range $[0, 1)$.
- Implement a method **forward(input_tensor)** which returns a tensor that serves as the **input_tensor** for the next layer. **input_tensor** is a matrix with **input_size** columns and **batch_size** rows. The **batch_size** represents the number of inputs processed simultaneously. The **output_size** is a parameter of the layer specifying the number of columns of the output.
- Add a setter and getter property **optimizer** which sets and returns the protected member **_optimizer** for this layer. Properties offer a pythonic way of realizing getters and setters. Please get familiar with this concept if you are not aware of it.
- Implement a method **backward(error_tensor)** which returns a tensor that serves as the **error_tensor** for the previous layer. Quick reminder: in the backward pass we are going in the other direction as in the forward pass.
Hint: if you discover that you need something here which is no longer available to you, think about storing it at the appropriate time.
- To be able to test the gradients with respect to the weights: The member for the weights and biases should be named **weights**. Additionally provide a property **gradient_weights** which returns the gradient with respect to the weights, after they have been calculated in the backward-pass. These properties are accessed by the unit tests and are therefore also important to pass the tests.
- Use the method **calculate_update(weight_tensor, gradient_tensor)** of your optimizer in your backward pass, in order to update your weights. **Don't perform an update if the optimizer is not set.**

4 Rectified Linear Unit

The Rectified Linear Unit is the standard activation function in Deep Learning nowadays. It has revolutionized Neural Networks because it reduces the effect of the “vanishing gradient” problem.

Task:

Implement a class **ReLU** in the file “ReLU.py” in folder “Layers”. This class also has to provide the methods **forward(input_tensor)** and **backward(error_tensor)**.

- Write a constructor for this class, receiving no arguments. The ReLU does not have trainable parameters, so you don’t have to change the inherited member trainable.
- Implement a method **forward(input_tensor)** which returns a tensor that serves as the **input_tensor** for the next layer.
- Implement a method **backward(error_tensor)** which returns a tensor that serves as the **error_tensor** for the previous layer.
Hint: the same hint as before applies.

5 SoftMax Layer

The SoftMax activation function is used to transform the logits (the output of the network) into a probability distribution. Therefore, SoftMax is typically used for classification tasks.

Task:

Implement a class **SoftMax** in the file: “SoftMax.py” in folder “Layers”. This class also has to provide the methods **forward(input_tensor)** and **backward(error_tensor)**.

- Write a constructor for this class, receiving no arguments.
- Implement a method **forward(input_tensor)** which returns the estimated class probabilities for each row representing an element of the batch.
- Implement a method **backward(error_tensor)** which returns a tensor that serves as the **error_tensor** for the previous layer.
Hint: again the same hint as before applies.
- Remember: Loops are slow in Python. Use NumPy functions instead!

6 Cross Entropy Loss

The cross entropy Loss is often used in classification task, typically in conjunction with SoftMax (or Sigmoid).

Task:

Implement a class **CrossEntropyLoss** in the file: “Loss.py” in folder “Optimization”. When forward propagating we now additionally need the argument **label_tensor** for **forward(input_tensor, label_tensor)** and **backward(label_tensor)**. We don’t consider the loss function as a layer like the previous ones in our framework, thus it should not inherit the base layer.

- Write a constructor for this class, receiving no arguments.
- Implement a method **forward(input_tensor, label_tensor)** which computes the Loss value according the CrossEntropy Loss formula accumulated over the batch.
- Implement a method **backward(label_tensor)** which returns the **error_tensor** for the previous layer. The backpropagation starts here, hence no **error_tensor** is needed. Instead, we need the **label_tensor**.
Hint: the same hint as before applies.
- Remember: Loops are slow in Python. Use NumPy functions instead!

7 Neural Network Skeleton

The Neural Network defines the whole architecture by containing all its layers from the input to the loss layer. This Network manages the testing and the training, that means it calls all forward methods passing the data from the beginning to the end, as well as the optimization by calling all backward passes afterwards.

Task:

Implement a class **NeuralNetwork** in the file: “NeuralNetwork.py” in the same folder as “NeuralNetworkTests.py”.

- Implement five public members. An **optimizer** object received upon construction as the first argument. A list **loss** which will contain the loss value for each iteration after calling **train**. A list **layers** which will hold the architecture, a member **data_layer**, which will provide input data and labels and a member **loss_layer** referring to the special layer providing loss and prediction. You do not need to care for filling these members with actual values. They will be set within the unit tests.
- Implement a method **forward** using input from the **data_layer** and passing it through all layers of the network. Note that the **data_layer** provides an **input_tensor** and a **label_tensor** upon calling **next()** on it. The output of this function should be the output of the last layer (i.e. the loss layer) of the network.
- Implement a method **backward** starting from the **loss_layer** passing it the **label_tensor** for the current input and propagating it back through the network.
- Implement the method **append_layer(layer)**. If the layer is **trainable**, it makes a deep_copy of the neural network's **optimizer** and sets it for the **layer** by using its optimizer property. Both, trainable and non-trainable layers, are then appended to the list **layers**.
Note: We will implement optimizers that have an internal state in the upcoming exercises, which makes copying of the optimizer object necessary.
- Additionally implement a convenience method **train(iterations)**, which trains the network for **iterations** and stores the loss for each iteration.
- Finally implement a convenience method **test(input_tensor)** which propagates the **input_tensor** through the network and returns the prediction of the last layer. For classification tasks we typically query the probabilistic output of the SoftMax layer.